



Introduction to Event Management

Event Management & Kafka

Introduction to Event Management

Event Management is a crucial process in today's fast-paced digital world. It involves capturing, processing, and analyzing events to gain valuable insights and drive business decisions. In this card, we'll explore the fundamentals of Event Management and its importance in optimizing operations and enhancing customer experiences.

Benefits of Event Sourcing

- Scalability
- Auditability
- Flexibility
- Real-time Responsiveness
- Data Integration
- Enhanced Monitoring

Event Management

- **Apache Kafka Apache**
- **Pulsar**
- **RabbitMQ**
- **Amazon Simple Queue Service (SQS)**
- **Google Cloud Pub/Sub**
- **Microsoft Azure Service Bus**
- **IBM MQ**

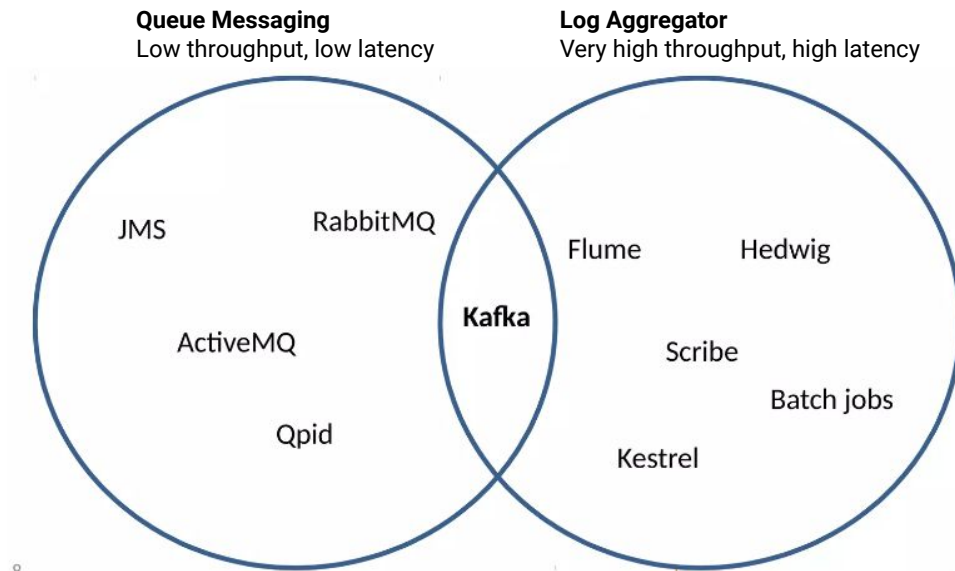
Event Management

Service	Description	Use Cases
Apache Kafka	Distributed streaming platform that focuses on high-throughput, fault-tolerant messaging, and real-time data processing.	Real-time data ingestion and processing, event streaming, log aggregation, messaging systems, and more.
RabbitMQ	Open-source message broker that implements the Advanced Message Queuing Protocol (AMQP) and provides reliable message delivery.	Enterprise messaging, task and workflow management, routing and filtering, and more.
Amazon Simple Queue Service (SQS)	Fully managed message queuing service that enables decoupling and scaling of distributed systems on the AWS platform.	Cloud-based messaging, decoupling of microservices, event-driven compute, and more.
Google Cloud Pub/Sub	Globally distributed messaging service that allows asynchronous communication between independently developed applications.	Real-time messaging, event-driven architectures, decoupling of microservices, and more.

Event Management

- **Apache Kafka:** Distributed streaming platform that focuses on high-throughput, fault-tolerant messaging, and real-time data processing.
- **Apache Pulsar:** Pub-sub messaging system that offers scalability, durability, and low-latency messaging, with support for both streaming and queuing models.
- **RabbitMQ:** Open-source message broker that implements the Advanced Message Queuing Protocol (AMQP) and provides reliable message delivery.
- **Amazon Simple Queue Service (SQS):** Fully managed message queuing service that enables decoupling and scaling of distributed systems on the AWS platform.
- **Google Cloud Pub/Sub:** Globally distributed messaging service that allows asynchronous communication between independently developed applications.
- **Microsoft Azure Service Bus:** Fully managed enterprise messaging service that provides asynchronous communication and decoupling for cloud-based applications.
- **IBM MQ:** Messaging middleware that facilitates communication and integration between applications, systems, and services across various platforms.

Message Queue Management Services



What is Kafka?

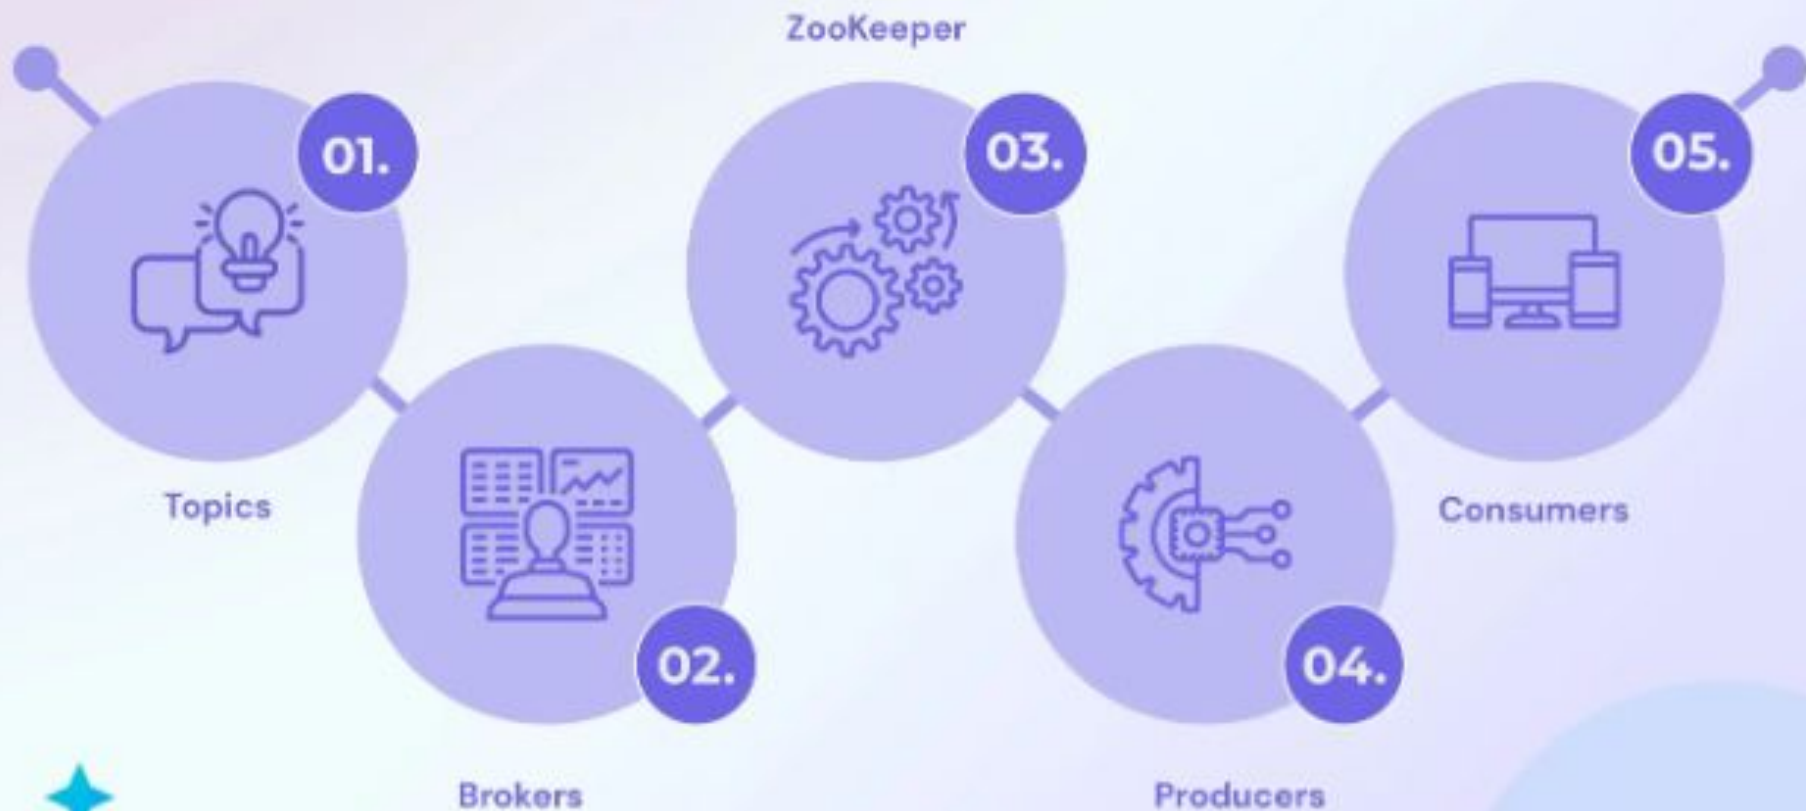
- Kafka is an open-source distributed event streaming platform
- It provides a highly scalable, fault-tolerant, and real-time data processing solution
- It allows organizations to efficiently capture, store, and process large volumes of streaming data, enabling seamless integration and real-time analytics



Concept of Kafka and its Specifications

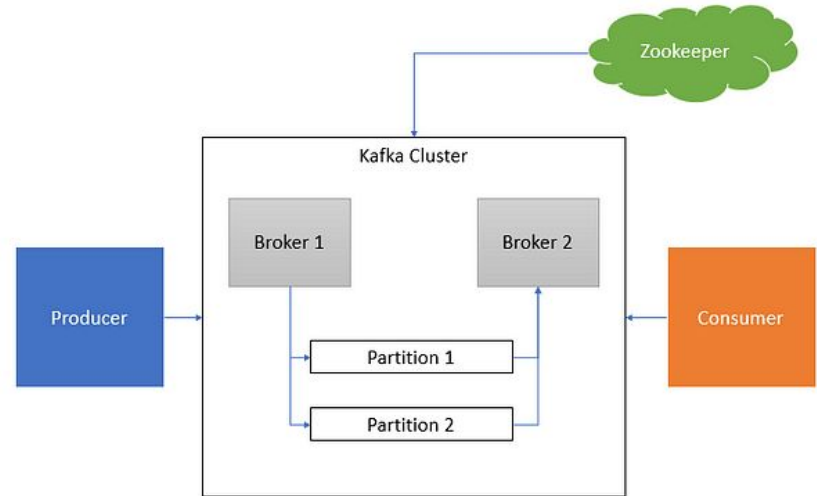


Kafka Clusters Architecture: 5 Major Components



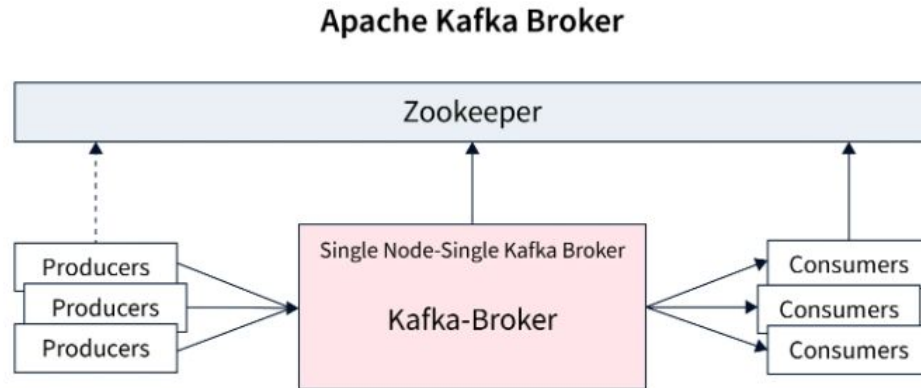
Kafka Architecture

- **Zookeeper:** Coordination service for discovery and registration of Kafka components
- **Kafka Cluster:** Set of servers called brokers that store and distribute messages
- **Producer:** Client that sends messages to Kafka topics
- **Partition:** Division of a Kafka topic to distribute messages across multiple brokers
- **Consumer:** Client that reads messages from one or more Kafka topics



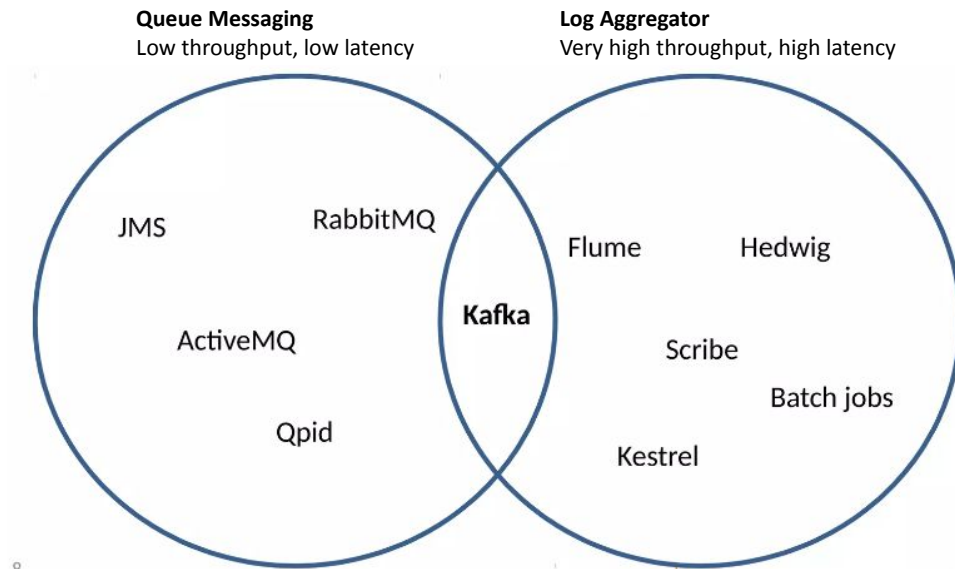
Kafka Architecture

- **Zookeeper**: Coordination service for discovery and registration of Kafka components
- **Kafka Cluster**: Set of servers called brokers that store and distribute messages
- **Producer**: Client that sends messages to Kafka topics
- **Partition**: Division of a Kafka topic to distribute messages across multiple brokers
- **Consumer**: Client that reads messages from one or more Kafka topics
- **Brokers** : Brokers are the servers that make up a Kafka cluster. They are responsible for storing and transmitting the messages within the group.



Comparing Kafka with Other Technologies

Comparison of Kafka with Other Brokers



Comparison of Kafka with Other Brokers

Feature	Kafka	ActiveMQ	RabbitMQ
Scalability	Designed for handling high volumes of data and supports horizontal scaling.	Can scale horizontally but may face performance issues with large data sets.	Can scale horizontally but may require additional configurations for high scalability.
Fault Tolerance	Replicates data across multiple brokers for automatic failover.	Provides fault tolerance but may require manual configuration for replication.	Provides fault tolerance through clustering and replication.
Real-time Processing	Enables low-latency message delivery for real-time data and stream processing.	Supports real-time processing but may have higher latency compared to Kafka.	Supports real-time processing but may have higher latency compared to Kafka.
Storage Efficiency	Stores data in a distributed and durable manner for efficient retention and retrieval.	Stores data but may have limitations on storage capacity and performance.	Stores data but may have limitations on storage capacity and performance.
Integration Capabilities	Integrates seamlessly with various data systems, databases, and frameworks.	Offers good integration capabilities but may require additional configurations.	Offers good integration capabilities but may require additional configurations.

Business Use Cases



LinkedIn

Activity stream, operational
metrics tracking, data bus



OVH

Anti-DDOS (OVH)



Netflix

Real-time tracking, real-time
processing (Netflix)



Twitter

Component of their real-time
architecture



Spotify

Log processing
(from 4am to 10am)



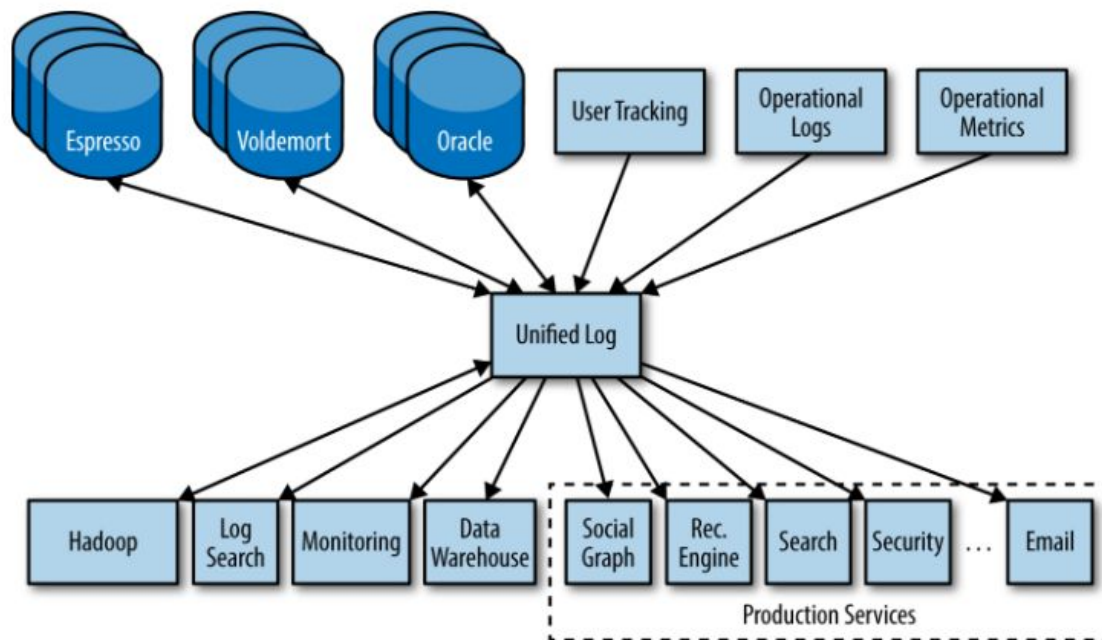
Mozilla

Metric management

Performance - LinkedIn

- 15 Brokers
- 15,500 partitions (replication factor 2)
- Input:
 - 400,000 messages/second
 - 70 MB/second
- Output:
 - 400 MB/second

Performance - LinkedIn



LinkedIn in



Performance - LinkedIn

Data Ingestion:

- Unified Log: Kafka is used to collect and store unified event logs from the system.
- Monitoring Data: Kafka is used to collect and store system monitoring data.

Data Processing:

- Espresso: Kafka is used to stream data to the Espresso data processing platform.
- Voldemort: Kafka is used to stream data to the Voldemort NoSQL data storage system.
- Hadoop: Kafka is used to stream data to the open-source Hadoop data processing platform.

Data Analysis:


- Log Search: Kafka is used to stream logs to the log search and analysis service.
- Recommendation Engine: Kafka is used to stream data to the recommendation engine for job postings and training.

Production Services:

- Production Services: Kafka is used to stream data to LinkedIn's production services.



Kafka - Spotify: Real-world Utilization



Kafka - Spotify: Real-world Utilization

Discover how Spotify leverages Kafka to handle millions of music streams and provide real-time recommendations to users, ensuring a seamless music streaming experience.

Example of Data:

Spotify's data includes information about user preferences, song metadata, and streaming activity.

Architecture:

Spotify's architecture involves multiple Kafka clusters for data ingestion, streaming, and processing. The data flows through various components, such as producers, topics, brokers, and consumers, enabling real-time data processing and recommendation generation.



Performance - Spotify

Kafka producers

10:22 UTC | 11:22 STO | 05:22 ASH | 02:22 SJC |

Page 1

Production Only

A

ash

accalia	afion	agnes	ahava	ainara	akana	ajesli	alma	ambika	anamika	andromache	anemone
UA 4451	UA 9853	UA 4875	UA 7205	UA 8322	UA 4283	UA 8018	UA 7686	UA 5683	UA 124	UA 7802	UA 7739
anissa	ann	annamiae	araluen	araminta	aranba	aretha	ash1-link-a1	ash1-link-a2	ash1-link-a3	ash1-link-a4	ash1-user/2-b1
UA 8123	UA 9056	UA 111	UA 5173	UA 7415	UA 5296	UA 4787	UA 6670	UA 6876	UA 8876	UA 20	UA 3028
ash1-user/2-b2	aurora	avery	bansali	barbro	berdine	bernadine	bhavya	bracha	bronnen	cameo	casandra
UA 2684	UA 5651	UA 7586	UA 7035	UA 2910	UA 4413	UA 4768	UA 4591	UA 3524	UA 7854	UA 7843	UA 7751
clauvery	chaitra	claudine	clementine	clermie	cleva	consuelo	cordelia	corinne	cyrana	daryl	dayana
UA 7307	UA 199	UA 4804	UA 6799	UA 7410	UA 5168	UA 10576	UA 10561	UA 8752	UA 7030	UA 8686	UA 13706
debbie	deborah	dietlinde	drisana	erica	estelle	fallon	felice	frankie	fumiko	gladys	gypsey
UA 15011	UA 7628	UA 6620	UA 5878	UA 2599	UA 2857	UA 6405	UA 7411	UA 8551	UA 102	UA 6622	UA 4630
halia	hanane	helvetia	herinda	iliaspes	iria	kajal	kenyatta	kismet	laurinda	lotta	lysaundra
UA 4786	UA 9414	UA 35	UA 8704	UA 5427	UA 4578	UA 5575	UA 6102	UA 8267	UA 10179	UA 4877	UA 8685
nediva	neeharika	nieves	pauline	rishbha	rosevear	samatha	samicah	sampriit	shradhdha	shulamit	stacia
UA 5043	UA 6362	UA 8065	UA 90	UA 125	UA 8954	UA 9883	UA 7999	UA 122	UA 6031	UA 3660	UA 8483
subhadra	surupa	tathra	velika								
UA 130	UA 127	UA 4883	UA 10								

ash1

ash1-link-a5	ash1-link-a6	ash1-link-a7	ash1-link-a8	ash1-notifications-a1	ash1-notifications-a2	ash1-notifications-a3	ash1-notifications-a4	ash1-notifications-a5	ash1-notifications-a6	ash1-notifications-a7
UA 38	UA 45	UA 35	UA 38	UA 0	UA 338	UA 288	UA 3	UA 3	UA 3	UA 3

ash2

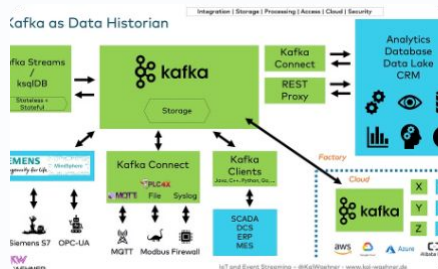
ash2-accesspoint-a1	ash2-accesspoint-a10	ash2-accesspoint-a11	ash2-accesspoint-a12	ash2-accesspoint-a13	ash2-accesspoint-a14	ash2-accesspoint-a15	ash2-accesspoint-a16	ash2-accesspoint-a17	ash2-accesspoint-a18	ash2-accesspoint-a19	ash2-accesspoint-a20
UA 13314	UA 14894	UA 6705	UA 9733	UA 11656	UA 10413	UA 11364	UA 14366	UA 9159	UA 9595	UA 13071	UA 8261
ash2-accesspoint-a21	ash2-accesspoint-a22	ash2-accesspoint-a23	ash2-accesspoint-a24	ash2-accesspoint-a25	ash2-accesspoint-a26	ash2-accesspoint-a27	ash2-accesspoint-a28	ash2-accesspoint-a29	ash2-accesspoint-a30	ash2-accesspoint-a31	ash2-accesspoint-a32
UA 8072	UA 7497	UA 14035	UA 14613	UA 8788	UA 14191	UA 8116	UA 14744	UA 16664	UA 8205	UA 14091	UA 11266
ash2-accesspoint-a33	ash2-accesspoint-a34	ash2-accesspoint-a35	ash2-accesspoint-a36	ash2-accesspoint-a37	ash2-accesspoint-a38	ash2-accesspoint-a39	ash2-accesspoint-a40	ash2-accesspoint-a41	ash2-accesspoint-a42	ash2-accesspoint-a43	ash2-accesspoint-a44
UA 9460	UA 12163	UA 8	UA 8999	UA 10073	UA 10418	UA 10981	UA 9089	UA 4151	UA 2884	UA 2823	UA 2465
ash2-user/2-a1	ash2-user/2-a2	ash2-user/2-a3	ash2-user/2-a4								
UA 1	UA 0	UA 2	UA 4								

Kafka - Spotify



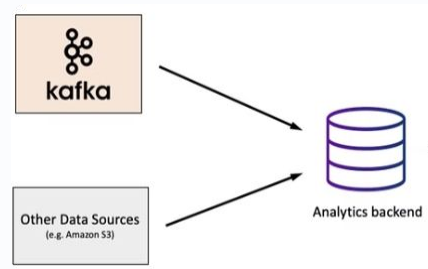
Music Streaming Integration

Kafka is integrated with Spotify for efficient music streaming and real-time data processing.



Real-time Data Processing

Spotify utilizes Kafka for real-time data processing and low-latency message delivery.



High Volume Event Handling

Kafka provides solutions for efficiently handling high volumes of event data for Spotify's platform.

Addressing Enterprise Challenges with Kafka



Kafka facilitates smooth integration with legacy systems, addressing compatibility challenges faced by enterprises.

Kafka offers solutions for efficiently handling high volumes of event data, a common challenge for enterprises.

Addressing Enterprise Challenges with Kafka



Kafka facilitates smooth integration with legacy systems, addressing compatibility challenges faced by enterprises.

Kafka offers solutions for efficiently handling high volumes of event data, a common challenge for enterprises.

Addressing Enterprise Challenges with Kafka



Kafka facilitates smooth integration with legacy systems, addressing compatibility challenges faced by enterprises.

Kafka offers solutions for efficiently handling high volumes of event data, a common challenge for enterprises.

Addressing Enterprise Challenges with Kafka




Kafka facilitates smooth integration with legacy systems, addressing compatibility challenges faced by enterprises.


Kafka offers solutions for efficiently handling high volumes of event data, a common challenge for enterprises.

Kafka's concepts

Streaming Integration with Kafka Connect



KAFKA CONNECT + STREAMS



Concepts de Kafka

- **Overview**

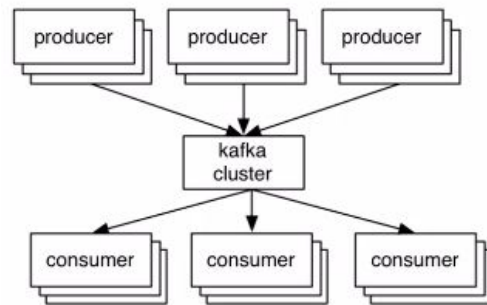
- Producers publish data to brokers.
- Consumers subscribe and retrieve data from brokers.
- All services are distributed.
- Data is stored in topics.

- **Details**

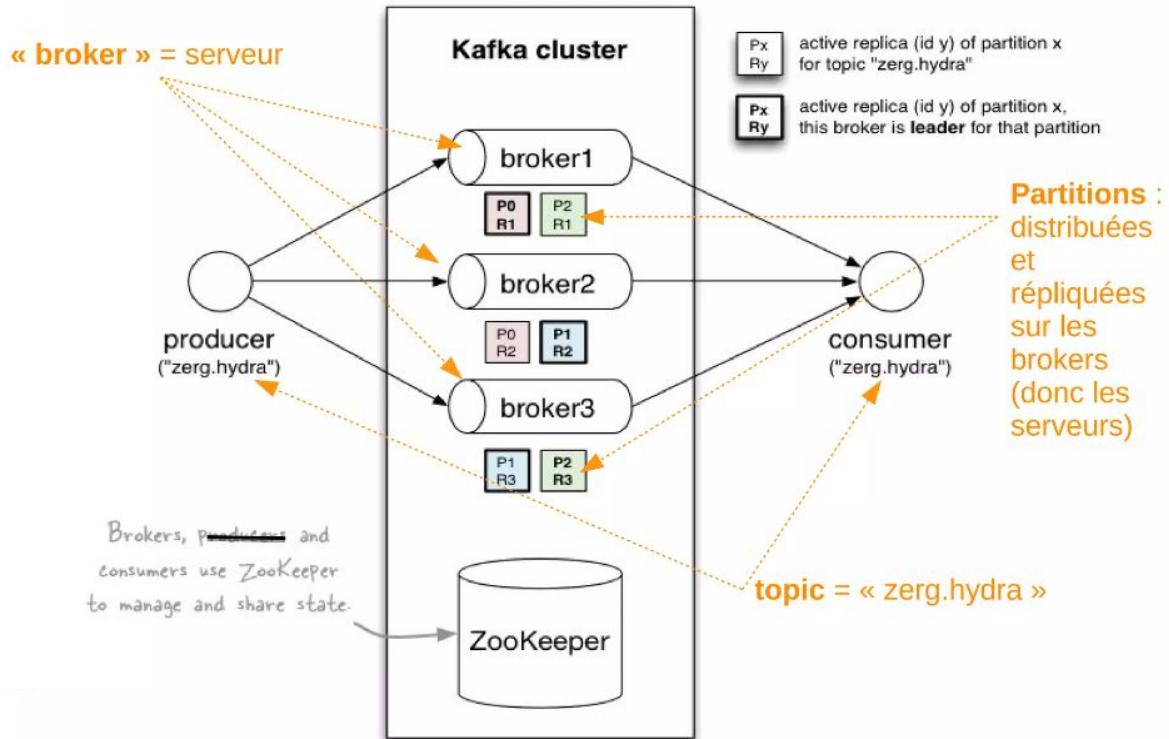
- Topics are divided into partitions and replicated to ensure high availability and scalability.
- Kafka is not a database system and does not handle SQL queries.

- **Advantages**

- High availability and scalability
- Low latency
- High throughput
- Asynchronous messaging
- Real-time streaming

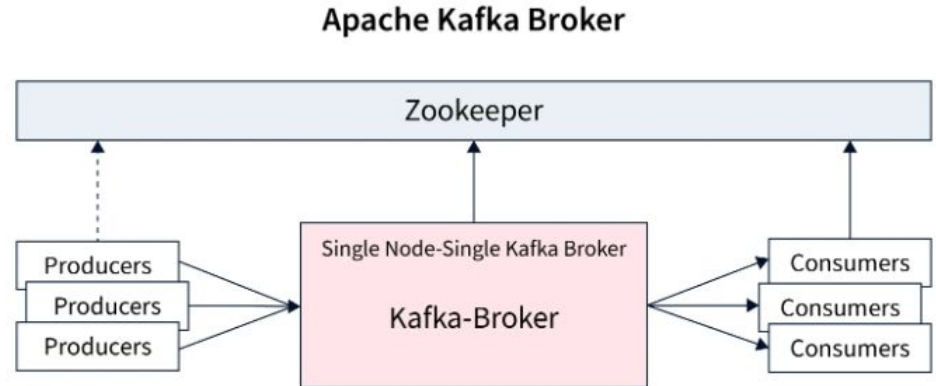


Concepts de Kafka - Overview



How Kafka Works

1. Producers send messages to Kafka topics.
2. Messages are stored on multiple brokers for fault tolerance.
3. Consumers read messages from brokers in order or disorder.
4. Consumers can subscribe to a single partition or a group of partitions.



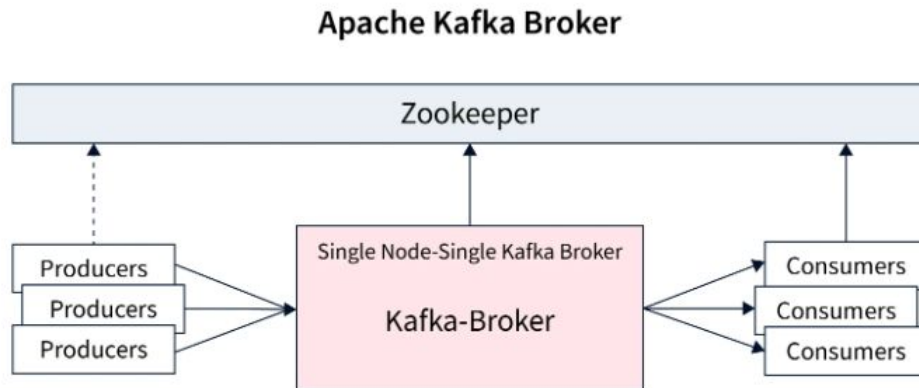
How to Start Kafka Broker?

```
./bin/zookeeper-server-start.sh config/zookeeper.properties
```

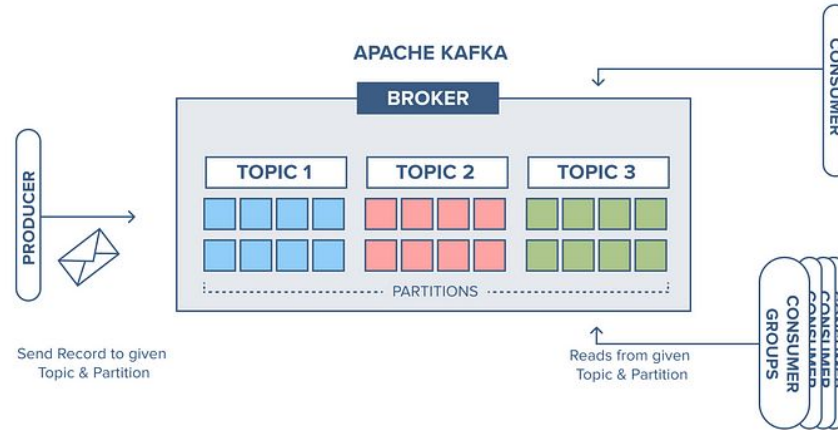
```
./bin/kafka-server-start.sh config/server.properties  
kafka-server-start.sh script
```

```
$ ./bin/kafka-server-start.sh
```

```
USAGE: ./bin/kafka-server-start.sh [-daemon] server.properties [--override property=value]*
```

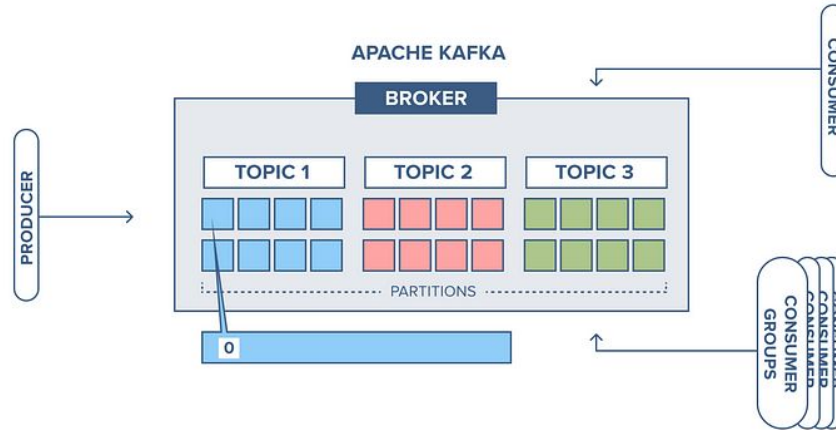


How Kafka Works



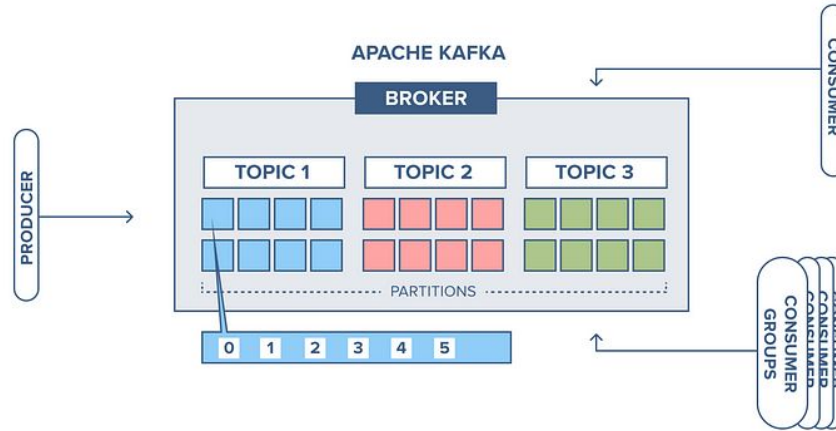
A broker with three topics, where each topic has 8 partitions.

How Kafka Works



The producer sends a record to partition 1 in topic 1 and since the partition is empty the record ends up at offset 0.

How Kafka Works



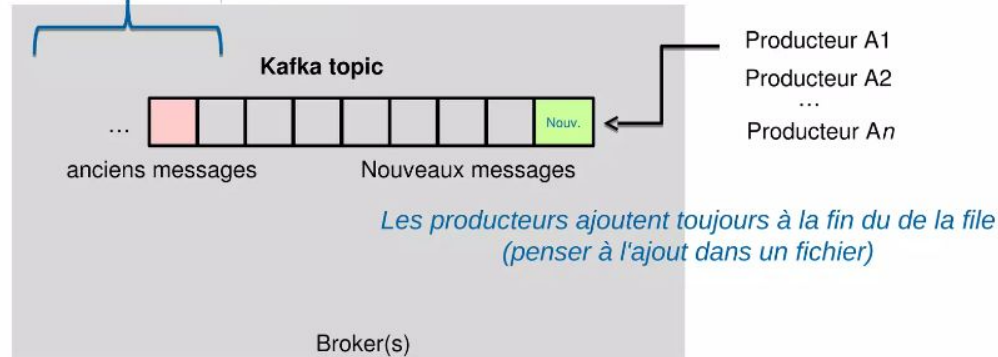
Next record is added to partition 1 will and up at offset 1, and the next record at offset 2 and so on.

Concepts de Kafka - Topic

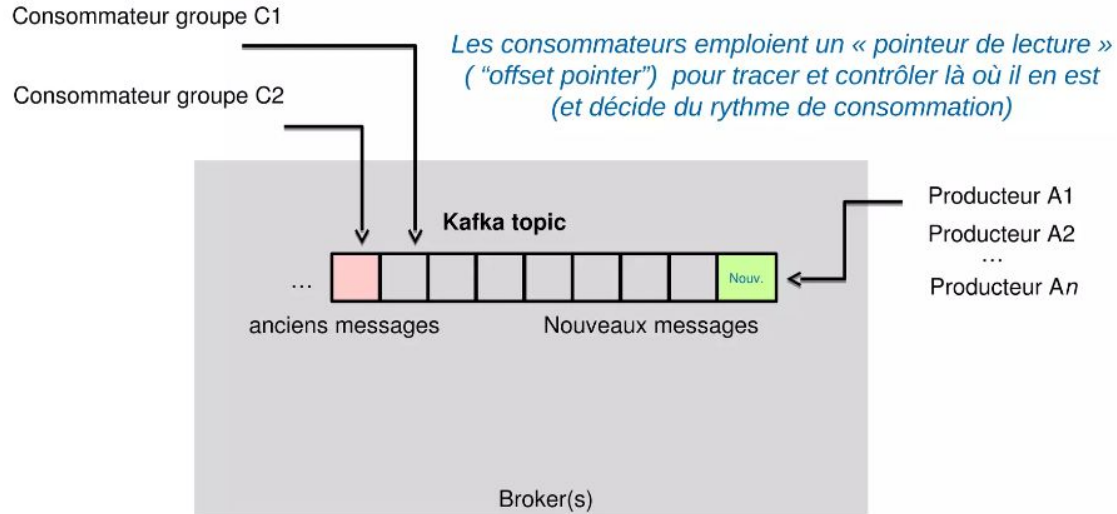
Topic: correspond au nom du flux sur lequel les messages vont être publiés

- Par exemple : “zerg.hydra”

Kafka élague depuis la “tête” en se basant sur l'âge ou la taille maximale or la « clé »



Concepts de Kafka - Topic

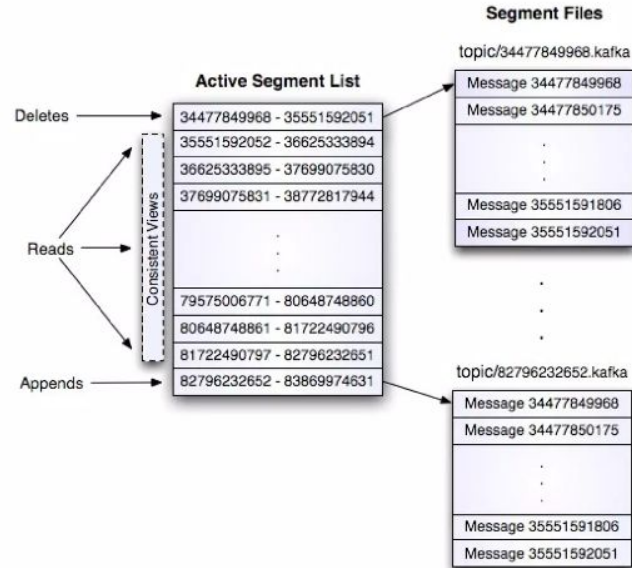



Concepts de Kafka - Message

Protocole léger

Traitement des messages par lot (Producteur & Consommateur)

Compression de bout en bout





Concepts de Kafka - Partition

Les partitions

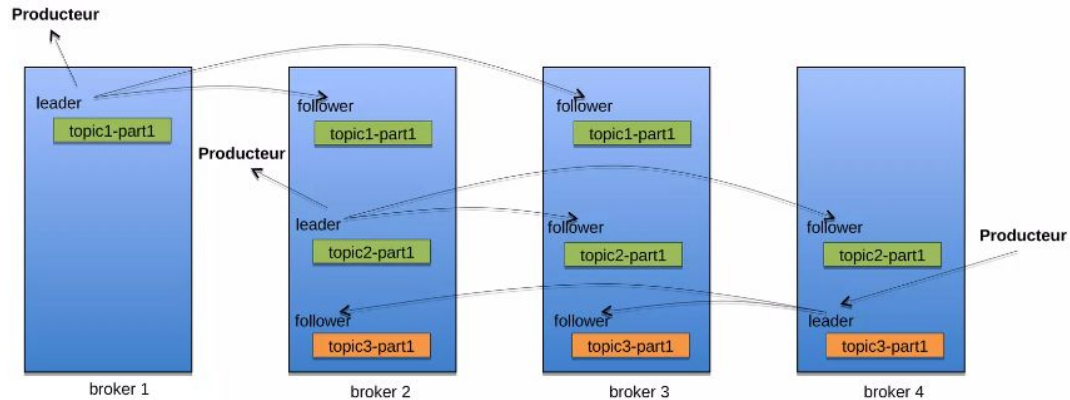

- Ordonnées
- Séquence immuable
- Le nombre de partitions détermine le nombre maximum de (groupes de) consommateurs

Anatomy of a Topic



Concepts de Kafka - Replicate

Flux de réplique



How Kafka Works

Steps to follow when setting up a connection and publishing a message/consuming a message.

1. First of all, we need to set up a secure connection. A TCP connection will be set up between the application and Apache Kafka.
2. In publisher: Publish a message to a partition on a topic.
3. In subscriber/consumer: Consume a message from a partition in a topic.

```

> import java.util.Properties
import kafkashaded.org.apache.kafka.clients.producer._
import org.apache.spark.sql.ForeachWriter

class KafkaSink(topic:String, servers:String) extends ForeachWriter[(String, String)] {
  val kafkaProperties = new Properties()
  kafkaProperties.put("bootstrap.servers", servers)
  kafkaProperties.put("key.serializer", "kafkashaded.org.apache.kafka.common.serialization.StringSerializer")
  kafkaProperties.put("value.serializer", "kafkashaded.org.apache.kafka.common.serialization.StringSerializer")
  val results = new scala.collection.mutable.HashMap[String, String]
  var producer: KafkaProducer[String, String] = _

  def open(partitionId: Long, version: Long): Boolean = {
    producer = new KafkaProducer(kafkaProperties)
    true
  }

  def process(value: (String, String)): Unit = {
    producer.send(new ProducerRecord(topic, value._1 + ":" + value._2))
  }

  def close(errorOrNull: Throwable): Unit = {
    producer.close()
  }
}

```

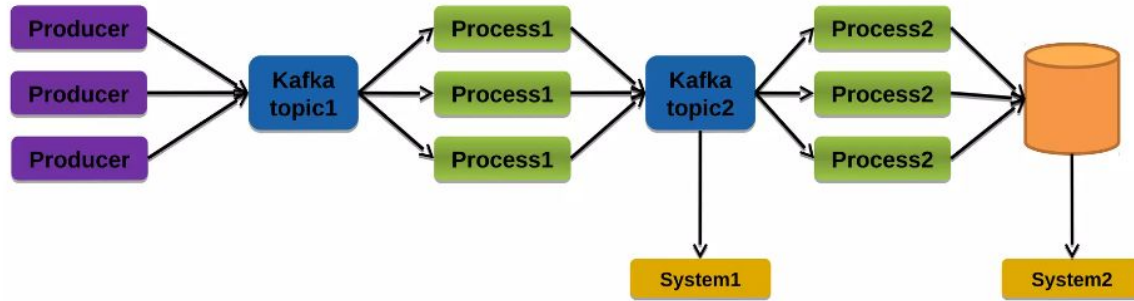
```

import java.util.Properties
import kafkashaded.org.apache.kafka.clients.producer._
import org.apache.spark.sql.ForeachWriter
defined class KafkaSink

```

Enhancing Event Scheduling with Kafka

Ecosystem





Apache Kafka adoption spans companies across industries.

Suite - Spring Batch Kafka
